



CHAPTER 14

BY :

ETISHA JAIN

VISHU GOYAL

LINUX 14.1 Understanding Disk Layout

a Disk have a partition like a pizza and also mounted

- **disk name = /dev/nvme0n1** It is specific type of advanced ssd
- **disk name = /dev/sda** It is the first SCSI disk of the system
- **disk name = /dev/vda** In KVM virtual machine

Partition depend on the disk name
like **sda1, sda2, sda3** and so on....

There are two option :

- your system is a **BIOS system**, BIOS comes from the original PC specification back from 1981, and it was invented to deal with systems of those days, In those days if system have 5MiB space then the computer is big so the BIOS is kind of outdated
- **UEFI -> Universal Extended Firmware Interface**

Different methods to address your disk

- **BIOS -> MBR** (master boot record) **64Byte** partition and 4 partition only **limitation**
- **so** we will use **logical partition extended** in the **4th partition** like **sda4**
- **by using logical partition** we can also make more partition in a disk like 5, 6, 7, 8 and so on..
- **UEFI -> GPT** (GUID Based Partition Table) **128 partitions**

- **lsblk** -> list block devices It shows all the disk present in our system
- **Sr0** is the SCSI CD_ROM drive
- **loop0** this is a block device which is used to mount the iso file
- some are **nvme0n1, nvme0n2, nvme0n3** and so on...

- **parted /dev/nvme0n1**
- **print**
- **quit**
- **cd /dev**
- **ls -l nvme***
- **cat /proc/partitions**

LINUX 14.2 Understanding Linux Storage Options

- **Partition** : The classical solution, use in all cases
 - Use to allocate dedicated storage to specific types of data
- **LVM logical Volumes**
 - use at default installation of RHEL
 - Adds flexibility to storage (resize, snapshots and more)
- **Stratis**
 - New feature in RHEL8
 - Next generation volume managing file system that uses thin provisioning by default
 - Implemented in user space, which makes API Access possible which means it can be addresses from different applications, which is a perfect solution in an environment where clouds or virtualization or containers are used
- **Virtual Data Optimizer**
 - Focused on storing files in the most efficient ways
 - Managing deduplicated and compressed storage pools

LINUX 14.3 Understanding GPT and MBR Partitions

difference between them:

- **Master Boot Record MBR** is a part of the 1981 PC specification
- 512 bytes to store boot information
- 64 bytes to store partitions
- place for 4 partitions only with a max. size of 2 TiB
- to use more partitions, extended and logical partitions must be used.
- **GUID Partition Table** is a newer partition table 2010
- More space to store partitions
- Used to overcome MBR limitations
- 128 partitions max
- by default UEFI server is used in GUID

LINUX 14.4 Creating Partitions with parted

- While **creating a partition**, you do not automatically create a file system
- The **parted file system** attribute only write some unimportant file system metadata
- In RHEL 8, **parted** is default utility
- Alternatively, use **fdisk** to work with **MBR** and **gdisk** to use **GUID** partitions

Steps to create a partitions

- **parted /dev/sdb**
- **print** will show if there is a current partition table
- **mklabel msdos | gpt** mklabel is used to set a partitions type and use msdos to set MBR or gpt to create a GUID partition table
- **mkpart part-type name fd-type start end**
 - **part-type** : applies to MBR only and sets primary, logical, or extended partition
 - **name** : arbitrary name, required for GPT
 - **fs-type** : does not modify the filesystem, but sets some irrelevant file system dependent metadata.
 - **start end** : specify start and end counting from the beginning of the disk.
- **for Instance** : mkpart primary 1024MiB 2048MiB.
- Alternatively, use **mkpart** in interactive mode
- **print** to verify creation of the new partition
- **quit** to exit the parted shell
- **udevadm settle** to ensure that the new partition device is created
- **cat /proc/partitions** to verify the creation of the partition

Practical

- lsblk
- parted /dev/nvme0n2
- print
- mklabel gpt
- print
- mkpart one 1MiB 1024 MiB
- print
- quit
- udevadm settle (to update all device)
- cat /proc/partitions
- lsblk

LINUX 14.5 Creating MBR Partitions with fdisk

- **lsblk**
- In the previous lecture we have created a GUID partition in nvme0n1 so we can't create a MBR partition
- so in this lecture we use new partition, nvme0n3 to use MBR
- **fdisk /dev/nvme0n3**
- **m** for help to show the menu
- **n** for new partition
- **p** for primary partition
- **default number**
- **first sector** default
- **last sector** +1G
- **w** to save a partition

if three partitions are created then if you want to create more partitions then you will use extended option

- **e** for extended partition
- **number**
- **first sector** default
- **last sector** default
- again same procedure.....
- **w** to save it.
- sometime you will get an error that is **partition is busy**
- then
- **partprobe** it will update all partitions

LINUX 14.6 Understanding File System Differences

- **XFS is the default file system**
 - Fast and scalable
 - Use CoW copy on write to guarantee data integrity
 - Size can be increased, not decreased
- **Ext4 was default in RHEL 6 and is still used**
 - Backward compatible to Ext2
 - Size can be increased and decreased
 - Uses Journal to guarantee data integrity
- **Other file system are available but less common**
- **Btrfs** was the promise of a new next generation file system
 - and RedHat experimented in RedHat 7
 - RedHat has decided not to move forward with Btrfs
 - RHEL 8 is using Stratis as an alternative for Btrfs

LINUX 14.7 Making and Mounting File Systems

- making a file system
- **mkfs.xfs** creates an XFS file system
- **mkfs.ext4** creates an Ext4 file system
- Use **mkfs.[Tab][Tab]** to show a list of available file systems.
- Do NOT use mkfs as it will create an Ext2 file system!
- After making the file system, you can mount it in runtime using the mount command.
- Use **umount** before disconnecting a device

Practical

- **lsblk**
- **mkfs.xfs /dev/nvme0n3p1** device which is not mount till yet
- **mount /dev/nvme0n3p1 /mnt**
- **mount**
- **mount | grep '^/'**
- **cd /mnt**
- **ls**
- **cp /etc/hosts .**
- **ls**
- **umount /dev/nvme0n3p1**
- **ls /mnt**
- **cd**
- **ls /mnt**
- **umount /mnt**

Practical

- **mkfs.ext4 --help**
- **mkfs.ext4 /dev/nvme0n3p2**

- **mkfs [tab] [tab]**
- vfat is file system that offers windows compatibility

LINUX 14.8 Mounting Partitions through /etc/fstab

- **/etc/fstab** is the main configuration file to persistently mount partitions
- **/etc/fstab** content is used to generate systemd mounts by the **systemd-fstab-generator** utility
- to update **systemd**, make sure to use **systemctl daemon-reload** after editing **/etc/fstab**

Practical

• **/name of the device /directory name type of the file 0-> dump utility**
0-> to check

- **vim/etc/fstab**
- **/dev/nvme0n3p1 /xfs xfs defaults 0 0**
- **/dev/nvme0n3p2 /ext4 ext4 defaults 0 0**
- **systemctl daemon-reload**
- **mkdir /xfs /ext4**
- **mount -a**
- **mount|**

LINUX 14.9 Managing Persistent Naming Attributes

In datacenter environments, block device names may change. Different solutions exist for persistent naming

- **UUID:** a **UUID is automatically** generated for each device that contains a file system or anything similar
- **Label:** while creating the file system, the option **-L** can be used to set an arbitrary name that can be used for mounting the file system
- **Unique device** names are created in `/dev/disk`

Practical

- `lsblk`
- `fdisk /dev/nvme0n3`
- `p`
- `n` add a new 6th partition
- `default`
- `+2G`
- `w`
- `mkfs.xfs /dev/nvme0n3p5` now create a file system on both of them

- **mkfs.ext4 /dev/nvme0n3p6**
- **cd /**
- **mkdir /books /article**
- **vim /etc/fstab**
- **/dev/nvme0n3p5 /books xfs defaults 0 0**
- **/dev/nvme0n3p6 /articles ext4 defaults 0 0**
- **mount -a**
- **mount**

- **fdisk /dev/nvme0n3**
- **d**
- number **5**
- failed to remove device is busy
- **vim /etc/fstab**
- remove that line number for 5th partition or commented
- **reboot**
- press **e** to edit the boot prompt
- now remove **rhgb quiet**
- **press ctrl x**
- **by** default the logical partition starts from 5
- press root pwd
- press **ctrl D**
- **vim /etc/fstab**
- remove or comment that line
- **lsblk**
- **blkid**

Set the Label

- `tune2fs --help`
 - `xfs_admin --help`
 - `tune2fs -L articles /dev/nvme0n3p5`
 - `blkid`
 - `vim /etc/fstab`
 - now remove the device name with label
 - **`LABEL=articles`**
 - `mount | grep article`
 - `mount -a`
 - `mount | grep article`
 - `vim /etc/fstable`
 - we can also replace the label with the UUID
- but the uuid is quite unreadable**

- `cd /dev/disk`
- `ls -l`
- There are 6 different styles to give the name of the device
- `id`
- `label`
- `partlabel`
- `partuuid`
- `path`
- `uuid`
- `ls by-path/`

Tips : focus on label and UUID

LINUX 14.10 Managing Systemd Mounts

- /etc/fstab mounts already are systemd mounts
- Mounts can be created using systemd .mount files
- Using .mount files allows you to be more specific in defining dependencies
- Use `systemctl cat tmp.mount` for an example

Practical

- **`systemctl cat tmp.mount`**
- **there is two section**
- **UNIT**
- **MOUNT** this one is important one.
- **what**
- **where**
- **type**
- **options**

- **`mount -a | grep tmp`**
- **`systemctl status tmp.mount`**
- **`systemctl enable --now tmp.mount`**
- **`systemctl status tmp.mount`**
- **`mount | grep tmp`**
- **`vim /etc/fstab`**
- now i comment that lable article and want to do with systemd mounts
- `cp /usr/lib/systemd/system/tmp.mount /etc/systemd/system/articles.mount`
- if you have a subdirectory then use that data-article
- `vim /etc/systemd/system/articles.mount`

[Unit]

- Desc
- Doc
- Conf
- Before

[Mount]

- **What = LABEL=articles**
- **where =/articles**
- **type=ext4**
- **option=defaults**
- systemctl daemon-reload
- systemctl status articles.mount
- umount /articles
- systemctl status articles.mount
- systemctl enable --now articles.mount

LINUX 14.11 Managing XFS File Systems

- It is a default file system
- the **xfsdump** utility can be used for creating backups of XFS formatted device and considers specific XFS attributes
 - xfsdump only works on a complete XFS device
 - xfsdump can make full backups (-l 0) or different levels of incremental backups
 - **xfsdump -l 0 -f /backupfiles/data.xfsdump /data** creates a full backup of the contents of the /data directory
- The xfsrestore command is used to restore a backup that was made with xfsdump
 - **xfsrestore -f /backupfiles/data.xfsdump /data**
- The **xfsrepair** command can be manually started to repair broken XFS file systems

LINUX 14.12 Creating a Swap Partition

- Swap is RAM that is emulated on disk
- All Linux Systems should have atleast some swap
- The amount of swap depends on the use of the server
- Swap can be created on any block device, including swap files
- While creating swap with parted, set file system to linux-swap
- After creatign the swap partition, use mkswap to create the swap FS
- Activate using swapon
- if you will create a swap partition with the fdisk command then will use type 82

Practicle

- parted /dev/nvme0n2
- print
- mkpart|
- swap
- linux-swap type is important in case of swap partition
- 1GiB
- 2GiB
- print
- quit

- lsblk
- mkswap /dev/nvme0n2p2
- free -m
- swapon /dev/nvme0n2p2
- free -m
- if you want to be the mounting of swap to be persistent the use that systemd or fstab file
- vim /etc/fstab
- /dev/nvme0n2p2 swap swap defaults 0 0
- reboot

The background is a blue gradient with decorative white circuit-like lines in the corners. The lines consist of straight segments and small circles, resembling a stylized electronic circuit.

THANK YOU